

Shakespeare AI: YouTube Presentation Script

A 2-Minute NotebookLM Podcast Screenplay

1. Hook & The Problem (0:00 - 0:15)

Host A So, we're looking at a classic dilemma. How do you expose a 400-year-old literary archive to the web without getting hacked by prompt injections?

Host B Exactly. Standard stateless servers bleed session logs and crash on mobile viewports. We needed a secure, intelligent playground.

2. Why Agents? (0:15 - 0:35)

Host A That's why we went agentic. Monolithic LLMs try to do everything and fail. Here, specialized subagents handle individual tasks.

Host B Right. One agent fetches local plays, another queries modern news, and a third philosophizes—each fully constrained by its own system instructions.

3. Architecture & Safety (0:35 - 0:55)

Host A The backend runs on the Agent Development Kit (ADK). A central router classifies your intent against a strict Pydantic schema.

Host B And security is baked in: the client sends a unique session ID header to isolate log files, while a pre-flight guardrail agent immediately blocks malicious injections.

4. The Build & Antigravity (0:55 - 1:15)

Host A For the build, we pair-programmed with Antigravity—a developer agent that optimized mobile momentum-scrolling and set up the Cloud Run pipeline.

Host B It even ran a custom BFS flood-fill script to strip background noise from the bouncing cybernetic Shakespeare visual companion.

5. The Live Demo (1:15 - 1:45)

Host A Look at this in action. The scholarship agent pulls real-time news, while the advice agent matches geopolitical instability to Julius Caesar.

Host B It's a stark reminder that we are doomed to repeat history unless we study how its patterns are here in the present. And it runs perfectly on mobile!

6. Outro & Call to Action (1:45 - 2:00)

Host A Fully open-source, built with the ADK, and safely hosted on Google Cloud Run.

Host B Clone the repo, drop in your API key, and stop repeating history. Link in the description!